

Computer Science 3A - CSC3A10

Lecture 10b: Search Trees

Academy of Computer Science and Software Engineering
University of Johannesburg



1 (2,4) Trees


- Multi-Way Search Tree
- Multi-Way Inorder Traversal
- Multi-Way Searching
- (2,4) Trees
- Height of a (2,4) Tree
- (2,4) Tree Insertion
- (2,4) Tree Analysis of Insertion
- (2,4) Tree Deletion
- (2,4) Tree Analysis of Deletion
- Implementing a Dictionary

2 Red-Black Trees

- From (2,4) to Red-Black Trees
- Red-Black Tree Properties
- Height of a Red-Black Tree
- Red-Black Tree Insertion

- Remediating a Double Red
- Restructuring
- Recolouring
- Red-Black Tree Analysis of Insertion
- Red-Black Tree Deletion
- Remediating a Double Black
- Red-Black Tree Reorganisation

(2,4) Trees



Multi-Way Search Tree

What if we could store two or more Entries together (such as the case in a Binary Search Tree) in hopes of gaining computational performance?

Enter Multi-Way Search Trees

Multi-Way Search Tree II

A multi-way search tree is an ordered tree such that:

- Each internal node has at least two children and stores $d - 1$ key-element items (k_i, o_i) , where d is the number of children
- For a node with children $v_1 v_2 \dots v_d$ storing keys $k_1 k_2 \dots k_{d-1}$
 - keys in the subtree of v_1 are less than k_1
 - keys in the subtree of v_i are between k_{i-1} and $k_i (i = 2, \dots, d - 1)$
 - keys in the subtree of v_d are greater than k_{d-1}
- The leaves store no items and serve as placeholders


Computer Science 3A - CSC3A10



Multi-Way Inorder Traversal

- We can extend the notion of inorder traversal from binary trees to multi-way search trees
- Namely, we visit item (k_i, o_i) of node v between the recursive traversals of the subtrees of v rooted at children v_i and v_{i+1}
- An inorder traversal of a multi-way search tree visits the keys in increasing order

Multi-Way Inorder Traversal II



Multi-Way Searching

Similar to search in a binary search tree


Each internal node with children $v_1 v_2 \dots v_d$ and keys $k_1 k_2 \dots k_{d-1}$

- $k = k_i (i = 1, \dots, d - 1)$: the search terminates successfully
- $k < k_1$: we continue the search in child v_1
- $k_{i-1} < k < k_i (i = 2, \dots, d - 1)$: we continue the search in child v_i
- $k > k_{d-1}$: we continue the search in child v_d

Reaching an external node terminates the search unsuccessfully

Multi-Way Searching

Example: search for 30




(2,4) Trees

A (2,4) tree (also called 2-4 tree or 2-3-4 tree) is a multi-way search with the following properties

- Node-Size Property: every internal node has at most four children
- Depth Property: all the external nodes have the same depth

Depending on the number of children, an internal node of a (2,4) tree is called a 2-node, 3-node or 4-node

(2,4) Trees II



Height of a (2,4) Tree


Theorem: A (2,4) tree storing n items has height $O(\log n)$

Proof:

- Let h be the height of a (2,4) tree with n items
- Since there are at least 2^i items at depth $i = 0, \dots, h-1$ and no items at depth h , we have $n \geq 1 + 2 + 4 + \dots + 2^{h-1} = 2^h - 1$
- Thus, $h \leq \log(n + 1)$

Searching in a (2,4) tree with n items takes $O(\log n)$ time

Height of a (2,4) Trees II




(2,4) Tree Insertion

We insert a new item (k, o) at the parent v of the leaf reached by searching for k

- We preserve the depth property but
- We may cause an overflow (i.e., node v may become a 5-node)

(2,4) Tree Insertion II

Example: inserting key 30 causes an overflow




Overflow and Split

We handle an overflow at a 5-node v with a split operation:

- let $v_1 \dots v_5$ be the children of v and $k_1 \dots k_4$ be the keys of v
- node v is replaced nodes v' and v''
 - v' is a 3-node with keys $k_1 k_2$ and children $v_1 v_2 v_3$
 - v'' is a 2-node with key k_4 and children $v_4 v_5$
- key k_3 is inserted into the parent u of v (a new root may be created)

The overflow may propagate to the parent node u

Academy of Computer Science and Software Engineering



(2,4) Insert Algorithm

```
1 Algorithm insert(k, o)
2 Step 1. We search for key k to locate the insertion node v
3 Step 2. We add the new entry (k, o) at node v
4 Step 3. while overflow(v)
5     if isRoot(v)
6         create a new empty root above v
7     v = split(v)
```

insert(k,o)

Analysis of Insertion

Let T be a (2,4) tree with n items


- Tree T has $O(\log n)$ height
- Step 1 takes $O(\log n)$ time because we visit $O(\log n)$ nodes
- Step 2 takes $O(1)$ time
- Step 3 takes $O(\log n)$ time because each split takes $O(1)$ time and we perform $O(\log n)$ splits

Thus, an insertion in a (2,4) tree takes $O(\log n)$ time

(2,4) Tree Deletion

- We reduce deletion of an entry to the case where the item is at the node with leaf children
- Otherwise, we replace the entry with its inorder successor (or, equivalently, with its inorder predecessor) and delete the latter entry

(2,4) Tree Deletion II



Underflow and Fusion


- Deleting an entry from a node v may cause an underflow, where node v becomes a 1-node with one child and no keys
- To handle an underflow at node v with parent u , we consider two cases:

Underflow and Fusion II

Case 1: the adjacent siblings of v are 2-nodes

- Fusion operation: we merge v with an adjacent sibling w and move an entry from u to the merged node v'
- After a fusion, the underflow may propagate to the parent u

Underflow and Fusion III




Underflow and Transfer

Case 2: an adjacent sibling w of v is a 3-node or a 4-node
Transfer operation:

- 1 we move a child of w to v
- 2 we move an item from u to v
- 3 we move an item from w to u

After a transfer, no underflow occurs

Academy of Computer Science and Software Engineering



Analysis of Deletion

Let T be a $(2,4)$ tree with n items (therefore Tree T has $O(\log n)$ height)

In a deletion operation

- We visit $O(\log n)$ nodes to locate the node from which to delete the entry
- We handle an underflow with a series of $O(\log n)$ fusions, followed by at most one transfer
- Each fusion and transfer takes $O(1)$ time


Thus, deleting an item from a $(2,4)$ tree takes $O(\log n)$ time

Implementing a Dictionary

Comparison of efficient dictionary implementations

	Search	Insert	Delete	Notes
Hash Table	1 expected	1 expected	1 expected	No ordered dictionary methods Simple to implement
Skip List	$\log n$ high prob.	$\log n$ high prob.	$\log n$ high prob.	Randomized insertion Simple to implement
(2,4) Tree	$\log n$ worst-case	$\log n$ worst-case	$\log n$ worst-case	Complex to implement

Red-Black Trees




From (2,4) to Red-Black Trees

A red-black tree is a representation of a $(2,4)$ tree by means of a binary tree whose nodes are coloured red or black


In comparison with its associated $(2,4)$ tree, a red-black tree has

- same logarithmic time performance
- simpler implementation with a single node type


From (2,4) to Red-Black Trees II



From (2,4) to Red-Black Trees III



From (2,4) to Red-Black Trees IV



Red-Black Tree Properties

A red-black tree can also be defined as a binary search tree that satisfies the following properties:

- **Root Property:** the root is black
- **External Property:** every leaf is black
- **Internal Property:** the children of a red node are black
- **Depth Property:** all the leaves have the same black depth

```
graph TD; 9((9)) --- 4((4)); 9 --- 15((15)); 4 --- 2((2)); 4 --- 6((6)); 6 --- 7((7)); 15 --- 12((12)); 15 --- 21((21)); 2 --- L1[ ]; 2 --- L2[ ]; 6 --- L3[ ]; 7 --- L4[ ]; 7 --- L5[ ]; 12 --- L6[ ]; 12 --- L7[ ]; 21 --- L8[ ]; 21 --- L9[ ];
```

Height of a Red-Black Tree

Theorem: A red-black tree storing n entries has height $O(\log n)$

Proof: The height of a red-black tree is at most twice the height of its associated (2,4) tree, which is $O(\log n)$


- The search algorithm for a binary search tree is the same as that for a binary search tree
- By the above theorem, searching in a red-black tree takes $O(\log n)$ time

Red-Black Tree Insertion

To perform operation $\text{insert}(k, o)$, we execute the insertion algorithm for binary search trees and color red the newly inserted node z unless it is the root

- We preserve the root, external, and depth properties
- If the parent v of z is black, we also preserve the internal property and we are done
- Else (v is red) we have a double red (i.e., a violation of the internal property), which requires a reorganization of the tree

Example where the insertion of 4 causes a double red:




Remedying a Double Red

Consider a double red with child z and parent v , and let w be the sibling of v

Case 1: w is black

- The double red is an incorrect replacement of a 4-node
- Restructuring: we change the 4-node replacement

Remedying a Double Red II




Remedying a Double Red III

Case 2: w is red

- The double red corresponds to an overflow
- Recolouring: we perform the equivalent of a split


Remedying a Double Red IV




Restructuring

- A restructuring remedies a child-parent double red when the parent red node has a black sibling
- It is equivalent to restoring the correct replacement of a 4-node
- The internal property is restored and the other properties are preserved

There are four restructuring configurations depending on whether the double red nodes are left or right children



- A recolouring remedies a child-parent double red when the parent red node has a red sibling
- The parent v and its sibling w become black and the grandparent u becomes red, unless it is the root
- It is equivalent to performing a split on a 5-node
- The double red violation may propagate to the grandparent u



Red-Black Insert Algorithm

```

1 Algorithm insert(k, o)
2 Step 1. We search for key k to locate the insertion node z
3
4 Step 2. We add the new entry (k, o) at node z and colour z red
5
6 Step 3. while doubleRed(z)
7     if isBlack(sibling(parent(z)))
8         z = restructure(z)
9     return
10    else { sibling(parent(z) is red }
11        z = recolour(z)

```

insert(k,o)


Analysis of Insertion

- Recall that a red-black tree has $O(\log n)$ height
- Step 1 takes $O(\log n)$ time because we visit $O(\log n)$ nodes
- Step 2 takes $O(1)$ time
- Step 3 takes $O(\log n)$ time because we perform
 - $O(\log n)$ recolorings, each taking $O(1)$ time, and
 - at most one restructuring taking $O(1)$ time
- Thus, an insertion in a red-black tree takes $O(\log n)$ time

Red-Black Tree Deletion

- To perform operation `remove(k)`, we first execute the deletion algorithm for binary search trees
- Let v be the internal node removed, w the external node removed, and r the sibling of w
 - If either v or r was red, we colour r black and we are done
 - Else (v and r were both black) we colour r double black, which is a violation of the internal property requiring a reorganization of the tree

Example where the deletion of 8 causes a double black:



Remedying a Double Black

The algorithm for remedying a double black node w with sibling y considers three cases:

- **Case 1:** y is black and has a red child: We perform a restructuring, equivalent to a transfer and we are done
- **Case 2:** y is black and its children are both black: We perform a recolouring, equivalent to a fusion, which may propagate up the double black violation
- **Case 3:** y is red: We perform an adjustment, equivalent to choosing a different representation of a 3-node, after which either Case 1 or Case 2 applies

Deletion in a red-black tree takes $O(\log n)$ time

Red-Black Tree Reorganisation

<i>Insertion</i>	<i>Remedy double red</i>	
Red-Black Tree Action	(2,4) Tree action	Result
Restructuring	Change of 4-node representation	Double red removed
Recolouring	Split	Double red removed or propagated up
<i>Deletion</i>	<i>Remedy double black</i>	
Red-Black Tree Action	(2,4) Tree action	Result
Restructuring	Transfer	Double black removed
Recolouring	Fusion	Double black removed or propagated up
Adjustment	Change of 3-node representation	Restructuring or re-colouring follows